

## Linux Scheduler Analysis

### Abstract:

For this paper we looked at three different Linux schedulers. These were the CFS scheduler, (used most by Linux OS) First in First Out Scheduler fifo, and Round Robin Scheduler RR. To assess the ability of each of these schedulers we ran them each with three different types of processes, one that was computation bound, one that was I/O bound and one that was a mix of both. Then we ran each of these three different processes at three different sizes a small, medium and large versions. We took the results of this process and analyzed the data looking at the amount of user time, cpu time, and the number of voluntary and involuntary context switches. Based on these data points we developed some conclusions about which schedulers performed best under which conditions.

### INTRODUCTION:

Coming up with efficient process scheduling policy is a task that computer engineers have been working on for years. Many different strategies have been developed over the past decades. A few of these scheduling policies have been shown through proof and analysis to be significantly better than others. For this project we have taken three such processes and analyzed their run time with three different types of processes using a large, medium and small amounts of simultaneous processes. This way we can see not only how these Schedulers compare with different types of processes, but also how the algorithms used in each scheduler scale as the size of data .

### METHOD:

The different types of processes we used are co, I/O bound, and a mix of computation and I/O bound. We used a program called pi which uses a randomized process that generates the number pi. This program is compute bound meaning that it uses up lots of CPU doing calculations. The program generates a more accurate number the higher the iterations are, so we used 100,000 iterations. This was our first developed benchmark.

We used a program that copies N bytes from an input file to an output file. The process writes 102400 bytes in blocks of 1024.


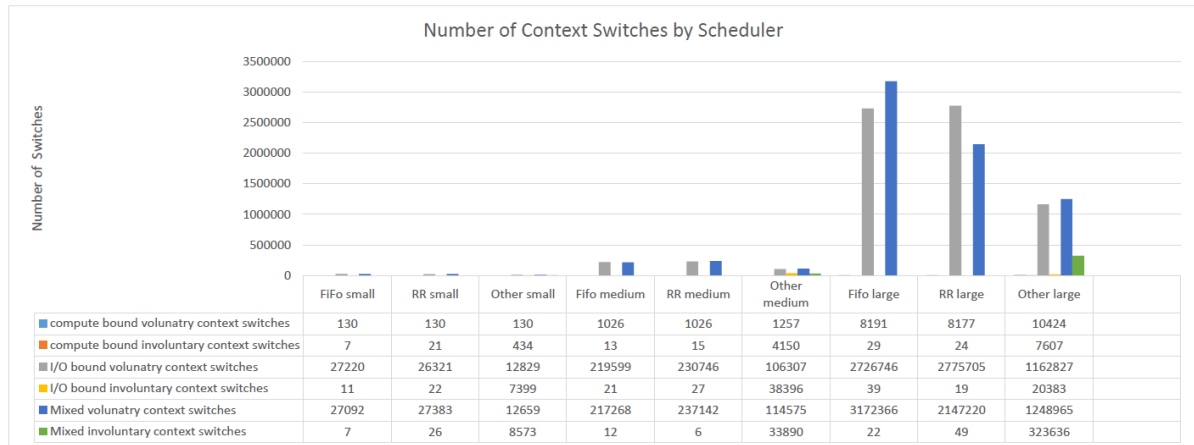
For the mixed I/O and compute bound we combined the pi process with the write to disk process, using the same parameters of iterations and writes as a benchmark.


We then developed size benchmarks as follows. The number of processes running simultaneously was considered to be small at 64, medium at 512 and large at 4096.

A bash script was written that would run all the test process in succession. We used a makefile to compile all of the test process. We set the -Wall and -Wextra error checking flags set to check for any errors. This bash script ran the pi process on each scheduler with the smallest simultaneous processes. Then it ran the I/O bound processes with the smallest simultaneous processes. It then proceeded to run the mixed I/O process with the smallest simultaneous processes. Once all the processes had finished running with the smallest number of simultaneous processes the script runs the medium (512) simultaneous processes in for each process and scheduler in the same order. Finally the largest (4096) simultaneous processes are run on all schedulers with all processes.

## RESULTS:

We have provided graphs which show a clear side by side comparison of the important factors. In the first graph these factors are the amount of user time and system time dedicated to each process by each scheduler across the small medium and large sets. In the second graph we show a comparison of involuntary and voluntary context switches for each process across all sets.





#### ANALYSIS:

We were expecting to see that the CFS scheduler performs the best since it is the Scheduler that is currently used by the Linux OS and stands to reason that they would implement the best scheduling policy.

From the results we can see important patterns that can help decide which scheduling algorithms perform best in each situation. A few general analysis first, as the number of simultaneous process being run increased the time dedicated to user time in compute bound increased, the amount of system time in I/O bound increased and the amount of both system and user time increased for the mixed compute and I/O bound process. This was a result that we were expecting and it was shown to be accurate. So we can determine that compute bound processes spend most of their time in the user space, while I/O bound processes spend most of their time in the kernel space.

With context switching we expected that the RR would have the most voluntary context switches followed by CFS and FIFO. For involuntary context switches we expected the FIFO to have the most followed by the CFS and the RR. Our prediction was inaccurate in both cases. The CFS has the most involuntary context switches for all size samples. Looking at the I/O bound context switches the RR and FIFO had very low involuntary context switches compared to the amount of voluntary context switches. For I/O bound medium size sample the RR and FIFO have 99.99% voluntary context switches. For the same process and sample size the CFS scheduler had 31% involuntary context switches. These finding surprised us since we expected to see higher involuntary context switches from the FIFO and lower involuntary context switches from the CFS. Also as the amount of simultaneous processes increased the number of voluntary and involuntary context switches increased. From this we can see that the more processes that are running in parallel the more context switches we can expect to see the CPU to make.

From this data we can start to understand which scheduler would be best suited for each different type of process and if this holds for all levels of concurrency. Looking at time the Other scheduler (Linux\_CFS) has the best performance by a very slight amount through all sizes of simultaneous runs. The difference between the three schedulers at the medium size

is smaller than at any other set size I would imagine this is do to the amount of involuntary context switching that occurs with the Other scheduler. Considering I/O bound processes, again we see the CFS scheduler consuming the least amount of time. They are all close to the same amount of time in the small size, the FIFO jumps up significantly at the medium size while the RR stays relatively close to the CFS. At the large size the FIFO is triple the time of the CFS and the RR is double the CFS. This significant gap as the size of the concurrent processes shows us that clearly the CFS is best suited for I/O bound processes. Looking at the mixed I/O and compute bound process again we see a very large jump in the system time for all three processes, 70 for FIFO, 60 for RR, and 30 for CFS.

## CONCLUSION

It is quite clear from the data provided that in terms of time the CFS scheduler out performs both the FIFO and the RR. We can develop a list of pro and cons from the data gathered in this report. CFS, pros: efficient run time allocation, this scheduler has the lowest time spent in system and user time of all the schedulers. Cons: high context switching and longer wait times. FIFO Pros: lower wait time, lower involuntary context switches. Cons: much higher system times necessary to complete I/O bound and mixed processes. RR, Pros: lowest wait times due to higher voluntary context switches. Cons: larger amount of time spent on system time in I/O bound processes. The use of CFS would be recommended for every day use of a computer. The RR scheduler would be good in a situation where you wanted to run a very large number of processes. The FIFO schedule would be good in a situation where you had a small number of Compute bound processes.

## APPENDIX A RAW DATA:

Building code...

1024+0 records in

1024+0 records out

1048576 bytes (1.0 MB) copied, 0.310337 s, 3.4 MB/s

Starting test runs...

Calculating pi over 100000 iterations using SCHED\_OTHER with 64 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_OTHER 6"

User time (seconds): 1.43

System time (seconds): 0.24

Percent of CPU this job got: 96%

Elapsed (wall clock) time (h:mm:ss or m:ss): 0:01.74

Average shared text size (kbytes): 0

Average unshared data size (kbytes): 0

Average stack size (kbytes): 0

Average total size (kbytes): 0

Maximum resident set size (kbytes): 4288

Average resident set size (kbytes): 0

Major (requiring I/O) page faults: 0

Minor (reclaiming a frame) page faults: 3022

Voluntary context switches: 130  
Involuntary context switches: 434  
Swaps: 0  
File system inputs: 0  
File system outputs: 520  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Calculating pi over 100000 iterations using SCHED\_FIFO with 64 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_FIFO 6"

User time (seconds): 1.69  
System time (seconds): 0.08  
Percent of CPU this job got: 191%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:00.93  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4300  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 3021  
Voluntary context switches: 130  
Involuntary context switches: 7  
Swaps: 0  
File system inputs: 0  
File system outputs: 528  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Calculating pi over 100000 iterations using SCHED\_RR with 64 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_RR 6"

User time (seconds): 1.60  
System time (seconds): 0.06  
Percent of CPU this job got: 181%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:00.92  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4192  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 3081  
Voluntary context switches: 130  
Involuntary context switches: 21  
Swaps: 0  
File system inputs: 0

File system outputs: 520  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_OTHER with 64 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 6  
SCHED\_OTHER"

User time (seconds): 0.01  
System time (seconds): 3.09  
Percent of CPU this job got: 46%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:06.74  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4240  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 2795  
Voluntary context switches: 12829  
Involuntary context switches: 7399  
Swaps: 0  
File system inputs: 0  
File system outputs: 207656  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_FIFO with 64 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 6 SCHED\_FIFO"

User time (seconds): 0.07  
System time (seconds): 8.50  
Percent of CPU this job got: 104%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:08.24  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4276  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 2734  
Voluntary context switches: 27220  
Involuntary context switches: 11  
Swaps: 0  
File system inputs: 0  
File system outputs: 393824

Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_RR with 64 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 6 SCHED\_RR"

User time (seconds): 0.05  
System time (seconds): 7.94  
Percent of CPU this job got: 104%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:07.65  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4268  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 2670  
Voluntary context switches: 26321  
Involuntary context switches: 22  
Swaps: 0  
File system inputs: 0  
File system outputs: 327904  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process

using SCHED\_OTHER with 64 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 6  
SCHED\_OTHER 100000"

User time (seconds): 1.26  
System time (seconds): 3.32  
Percent of CPU this job got: 57%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:07.92  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4216  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 3300  
Voluntary context switches: 12659  
Involuntary context switches: 8573  
Swaps: 0  
File system inputs: 0  
File system outputs: 217512  
Socket messages sent: 0

Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process  
using SCHED\_FIFO with 64 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 6  
SCHED\_FIFO 100000"

User time (seconds): 2.06  
System time (seconds): 8.17  
Percent of CPU this job got: 115%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:08.85  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4148  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 3300  
Voluntary context switches: 27092  
Involuntary context switches: 7  
Swaps: 0  
File system inputs: 0  
File system outputs: 414408  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process  
using SCHED\_RR with 64 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 6 SCHED\_RR  
100000"

User time (seconds): 1.93  
System time (seconds): 7.93  
Percent of CPU this job got: 113%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:08.69  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4216  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 3175  
Voluntary context switches: 27383  
Involuntary context switches: 26  
Swaps: 0  
File system inputs: 0  
File system outputs: 425760  
Socket messages sent: 0



Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Calculating pi over 100000 iterations using SCHED\_OTHER with 512 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_OTHER 9"  
User time (seconds): 12.38  
System time (seconds): 0.53  
Percent of CPU this job got: 93%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:13.79  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4276  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 21233  
Voluntary context switches: 1257  
Involuntary context switches: 4150  
Swaps: 0  
File system inputs: 0  
File system outputs: 4200  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Calculating pi over 100000 iterations using SCHED\_FIFO with 512 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_FIFO 9"  
User time (seconds): 10.05  
System time (seconds): 0.43  
Percent of CPU this job got: 190%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:05.50  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4216  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 19691  
Voluntary context switches: 1026  
Involuntary context switches: 13  
Swaps: 0  
File system inputs: 0  
File system outputs: 4184  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0

Page size (bytes): 4096  
Exit status: 0  
Calculating pi over 100000 iterations using SCHED\_RR with 512 simultaneous processes...  
Command being timed: "sudo ./pi-sched 100000 SCHED\_RR 9"  
User time (seconds): 12.04  
System time (seconds): 0.62  
Percent of CPU this job got: 190%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:06.65  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4280  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 21232  
Voluntary context switches: 1026  
Involuntary context switches: 15  
Swaps: 0  
File system inputs: 0  
File system outputs: 4200  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0  
Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_OTHER with 512 simultaneous processes...  
Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 9  
SCHED\_OTHER"  
User time (seconds): 0.11  
System time (seconds): 14.16  
Percent of CPU this job got: 45%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:31.07  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4248  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 17389  
Voluntary context switches: 106307  
Involuntary context switches: 38396  
Swaps: 0  
File system inputs: 0  
File system outputs: 857528  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_FIFO with 512 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 9 SCHED\_FIFO"

User time (seconds): 0.31

System time (seconds): 55.77

Percent of CPU this job got: 125%

Elapsed (wall clock) time (h:mm:ss or m:ss): 0:44.56

Average shared text size (kbytes): 0

Average unshared data size (kbytes): 0

Average stack size (kbytes): 0

Average total size (kbytes): 0

Maximum resident set size (kbytes): 4216

Average resident set size (kbytes): 0

Major (requiring I/O) page faults: 0

Minor (reclaiming a frame) page faults: 18925

Voluntary context switches: 219599

Involuntary context switches: 21

Swaps: 0

File system inputs: 0

File system outputs: 2509112

Socket messages sent: 0

Socket messages received: 0

Signals delivered: 0

Page size (bytes): 4096

Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_RR with 512 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 9 SCHED\_RR"

User time (seconds): 0.19

System time (seconds): 19.97

Percent of CPU this job got: 84%

Elapsed (wall clock) time (h:mm:ss or m:ss): 0:23.93

Average shared text size (kbytes): 0

Average unshared data size (kbytes): 0

Average stack size (kbytes): 0

Average total size (kbytes): 0

Maximum resident set size (kbytes): 4192

Average resident set size (kbytes): 0

Major (requiring I/O) page faults: 0

Minor (reclaiming a frame) page faults: 17385

Voluntary context switches: 230746

Involuntary context switches: 27

Swaps: 0

File system inputs: 0

File system outputs: 2112304

Socket messages sent: 0

Socket messages received: 0

Signals delivered: 0

Page size (bytes): 4096

Exit status: 0

Running mixed CPU/IO bound process

using SCHED\_OTHER with 512 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 9  
SCHED\_OTHER 100000"

User time (seconds): 2.71  
System time (seconds): 5.20  
Percent of CPU this job got: 35%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:22.47  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4216  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 22502  
Voluntary context switches: 114575  
Involuntary context switches: 33890  
Swaps: 0  
File system inputs: 0  
File system outputs: 919744  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process  
using SCHED\_FIFO with 512 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 9  
SCHED\_FIFO 100000"

User time (seconds): 3.94  
System time (seconds): 9.74  
Percent of CPU this job got: 50%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:26.96  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4200  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 22514  
Voluntary context switches: 217268  
Involuntary context switches: 12  
Swaps: 0  
File system inputs: 0  
File system outputs: 1350984  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process  
using SCHED\_RR with 512 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 9 SCHED\_RR 100000"

User time (seconds): 3.99  
System time (seconds): 10.64  
Percent of CPU this job got: 78%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:18.53  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4196  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 21485  
Voluntary context switches: 237142  
Involuntary context switches: 6  
Swaps: 0  
File system inputs: 0  
File system outputs: 1602984  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Calculating pi over 100000 iterations using SCHED\_OTHER with 4096 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_OTHER 12"

User time (seconds): 26.02  
System time (seconds): 2.31  
Percent of CPU this job got: 99%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:28.55  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4196  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 145907  
Voluntary context switches: 10424  
Involuntary context switches: 7607  
Swaps: 0  
File system inputs: 0  
File system outputs: 33536  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Calculating pi over 100000 iterations using SCHED\_FIFO with 4096 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_FIFO 12"

User time (seconds): 25.74  
System time (seconds): 1.75  
Percent of CPU this job got: 190%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:14.44  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4252  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 162298  
Voluntary context switches: 8191  
Involuntary context switches: 29  
Swaps: 0  
File system inputs: 0  
File system outputs: 33536  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Calculating pi over 100000 iterations using SCHED\_RR with 4096 simultaneous processes...

Command being timed: "sudo ./pi-sched 100000 SCHED\_RR 12"

User time (seconds): 24.81  
System time (seconds): 1.81  
Percent of CPU this job got: 190%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:13.99  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4196  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 150007  
Voluntary context switches: 8177  
Involuntary context switches: 24  
Swaps: 0  
File system inputs: 0  
File system outputs: 33536  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_OTHER with 4096 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 12

SCHED\_OTHER"

User time (seconds): 0.14  
System time (seconds): 34.61

Percent of CPU this job got: 50%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 1:09.13  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4276  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 152058  
Voluntary context switches: 1162827  
Involuntary context switches: 280383  
Swaps: 0  
File system inputs: 0  
File system outputs: 5091520  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_FIFO with 4096 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 12 SCHED\_FIFO"  
User time (seconds): 0.30  
System time (seconds): 102.78  
Percent of CPU this job got: 84%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 2:02.31  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4216  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 139803  
Voluntary context switches: 2726746  
Involuntary context switches: 39  
Swaps: 0  
File system inputs: 0  
File system outputs: 8569144  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Copying 102400 bytes in blocks of 1024 from rwinput to rwoutput  
using SCHED\_RR with 4096 simultaneous processes...

Command being timed: "sudo ./rw 102400 1024 rwinput rwoutput 12 SCHED\_RR"  
User time (seconds): 0.22  
System time (seconds): 77.87  
Percent of CPU this job got: 66%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 1:58.33

Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4252  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 143935  
Voluntary context switches: 2775705  
Involuntary context switches: 19  
Swaps: 0  
File system inputs: 0  
File system outputs: 7427072  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process

using SCHED\_OTHER with 4096 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 12

SCHED\_OTHER 100000"

User time (seconds): 21.32  
System time (seconds): 36.47  
Percent of CPU this job got: 55%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 1:44.46  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4200  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 176623  
Voluntary context switches: 1248965  
Involuntary context switches: 323636  
Swaps: 0  
File system inputs: 0  
File system outputs: 5203440  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process

using SCHED\_FIFO with 4096 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 12

SCHED\_FIFO 100000"

User time (seconds): 31.31  
System time (seconds): 84.80  
Percent of CPU this job got: 77%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 2:29.78



Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4188  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 184871  
Voluntary context switches: 3172366  
Involuntary context switches: 22  
Swaps: 0  
File system inputs: 0  
File system outputs: 9274960  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

Running mixed CPU/IO bound process  
using SCHED\_RR with 4096 simultaneous processes...

Command being timed: "sudo ./mixed 102400 1024 rwinput rwoutput 12

SCHED\_RR 100000"

User time (seconds): 30.94  
System time (seconds): 70.28  
Percent of CPU this job got: 62%  
Elapsed (wall clock) time (h:mm:ss or m:ss): 2:42.72  
Average shared text size (kbytes): 0  
Average unshared data size (kbytes): 0  
Average stack size (kbytes): 0  
Average total size (kbytes): 0  
Maximum resident set size (kbytes): 4196  
Average resident set size (kbytes): 0  
Major (requiring I/O) page faults: 0  
Minor (reclaiming a frame) page faults: 180756  
Voluntary context switches: 2147220  
Involuntary context switches: 49  
Swaps: 0  
File system inputs: 0  
File system outputs: 8176656  
Socket messages sent: 0  
Socket messages received: 0  
Signals delivered: 0  
Page size (bytes): 4096  
Exit status: 0

user@cu-cs-vm:~/Dropbox/3753/PA4/pa3\$

## APPENDIX B: ALL CODE

-mixed.c-

/\*

\* File: pi-sched.c

\* Author: Andy Sayler

```

* Project: CSCI 3753 Programming Assignment 3
* Create Date: 2012/03/07
* Modify Date: 2012/03/09
* Description:
*     This file contains a simple program for statistically
*     calculating pi using a specific scheduling policy.
*/

/*
* File: rw.c
* Author: Andy Sayler
* Project: CSCI 3753 Programming Assignment 3
* Create Date: 2012/03/19
* Modify Date: 2012/03/20
* Description: A small i/o bound program to copy N bytes from an input
*     file to an output file. May read the input file multiple
*     times if N is larger than the size of the input file.
*/

/* Include Flags */
#define _GNU_SOURCE

/* System Includes */
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <errno.h>
#include <sched.h>
#include <math.h>
#include <fcntl.h>
#include <string.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/wait.h>

#define DEFAULT_ITERATIONS 1000000
#define RADIUS (RAND_MAX / 2)
#define DEFAULT_PROCESSES 1
#define MAXFILENAMELENGTH 80
#define DEFAULT_INPUTFILENAME "rwinput"
#define DEFAULT_OUTPUTFILENAMEBASE "rwoutput"
#define DEFAULT_BLOCKSIZE 1024
#define DEFAULT_TRANSFERSIZE 1024*100
#define DEFAULT_PROCESSES 1

inline double dist(double x0, double y0, double x1, double y1){
    return sqrt(pow((x1-x0),2) + pow((y1-y0),2));
}

inline double zeroDist(double x, double y){
    return dist(0, 0, x, y);
}

```

```

int main(int argc, char* argv[]){

    int rv;
    int inputFD;
    int outputFD;
    int status;
    int policy;
    long processes;
    struct sched_param param;
    char inputFilename[MAXFILENAMELENGTH];
    char outputFilename[MAXFILENAMELENGTH];
    char outputFilenameBase[MAXFILENAMELENGTH];

    long i;
    long iterations;
    double x, y;
    double inCircle = 0.0;
    double inSquare = 0.0;
    double pCircle = 0.0;
    double piCalc = 0.0;

    ssize_t transfersize = 0;
    ssize_t blocksize = 0;
    char* transferBuffer = NULL;
    ssize_t buffersize;

    ssize_t bytesRead = 0;
    ssize_t totalBytesRead = 0;
    int totalReads = 0;
    ssize_t bytesWritten = 0;
    ssize_t totalBytesWritten = 0;
    int totalWrites = 0;
    int inputFileResets = 0;

    /* Process program arguments to select run-time parameters */
    /* Set supplied transfer size or default if not supplied */
    if(argc < 2){
        transfersize = DEFAULT_TRANSFERSIZE;
    }
    else{
        transfersize = atol(argv[1]);
        if(transfersize < 1){
            fprintf(stderr, "Bad transfersize value\n");
            exit(EXIT_FAILURE);
        }
    }
    /* Set supplied block size or default if not supplied */
    if(argc < 3){
        blocksize = DEFAULT_BLOCKSIZE;
    }
    else{

```

```

        blocksize = atol(argv[2]);
        if(blocksize < 1){
            fprintf(stderr, "Bad blocksize value\n");
            exit(EXIT_FAILURE);
        }
    }
    /* Set supplied input filename or default if not supplied */
    if(argc < 4){
        if(strlen(DEFAULT_INPUTFILENAME, MAXFILENAMELENGTH) >=
MAXFILENAMELENGTH){
            fprintf(stderr, "Default input filename too long\n");
            exit(EXIT_FAILURE);
        }
        strncpy(inputFilename, DEFAULT_INPUTFILENAME,
MAXFILENAMELENGTH);
    }
    else{
        if(strlen(argv[3], MAXFILENAMELENGTH) >= MAXFILENAMELENGTH){
            fprintf(stderr, "Input filename too long\n");
            exit(EXIT_FAILURE);
        }
        strncpy(inputFilename, argv[3], MAXFILENAMELENGTH);
    }
    /* Set supplied output filename base or default if not supplied */
    if(argc < 5){
        if(strlen(DEFAULT_OUTPUTFILENAMEBASE, MAXFILENAMELENGTH) >=
MAXFILENAMELENGTH){
            fprintf(stderr, "Default output filename base too long\n");
            exit(EXIT_FAILURE);
        }
        strncpy(outputFilenameBase, DEFAULT_OUTPUTFILENAMEBASE,
MAXFILENAMELENGTH);
    }
    else{
        if(strlen(argv[4], MAXFILENAMELENGTH) >= MAXFILENAMELENGTH){
            fprintf(stderr, "Output filename base is too long\n");
            exit(EXIT_FAILURE);
        }
        strncpy(outputFilenameBase, argv[4], MAXFILENAMELENGTH);
    }
    /* Set default policy if not supplied */
    if(argc < 7){
        policy = SCHED_OTHER;
    }
    /* Set policy if supplied */
    if(argc > 6){
        if(!strcmp(argv[6], "SCHED_OTHER")){
            policy = SCHED_OTHER;
        }
        else if(!strcmp(argv[6], "SCHED_FIFO")){
            policy = SCHED_FIFO;
        }
    }

```

```

        else if(!strcmp(argv[6], "SCHED_RR")){
            policy = SCHED_RR;
        }
        else{
            fprintf(stderr, "Unhanded scheduling policy\n");
            exit(EXIT_FAILURE);
        }
    }
    /* Set default iterations if not supplied */
    if(argc < 8){
        iterations = DEFAULT_ITERATIONS;
    }
    /* Set iterations if supplied */
    if(argc > 7){
        iterations = atol(argv[7]);
        if(iterations < 1){
            fprintf(stderr, "Bad iterations value\n");
            exit(EXIT_FAILURE);
        }
    }
    /* Set process to max prioty for given scheduler */
    param.sched_priority = sched_get_priority_max(policy);
    /* Set new scheduler policy */
    fprintf(stdout, "Current Scheduling Policy: %d\n", sched_getscheduler(0));
    fprintf(stdout, "Setting Scheduling Policy to: %d\n", policy);
    if(sched_setscheduler(0, policy, &param)){
        perror("Error setting scheduler policy");
        exit(EXIT_FAILURE);
    }
    fprintf(stdout, "New Scheduling Policy: %d\n", sched_getscheduler(0));

    /* Modification for forks */
    if(argc < 6){
        processes = DEFAULT_PROCESSES;
    }

    /* Modification for forks */
    if(argc > 5){
        processes = atol(argv[5]);
        if(processes < 1){
            fprintf(stderr, "Bad processes value\n");
            exit(EXIT_FAILURE);
        }
    }

    /* Modification for forks */
    int f;
    int g;
    for(f=0; f<processes; f++){
        if(fork() == 0){
            execve(argv[1],argv, NULL);
        }
    }

```

```

    }
    for(g=0; g<processes; g++){
        waitpid(0, &status, 0);
    }

/* Confirm blocksize is multiple of and less than transfersize*/
if(blocksize > transfersize){
    fprintf(stderr, "blocksize can not exceed transfersize\n");
    exit(EXIT_FAILURE);
}
if(transfersize % blocksize){
    fprintf(stderr, "blocksize must be multiple of transfersize\n");
    exit(EXIT_FAILURE);
}

/* Allocate buffer space */
bufferSize = blockSize;
if(!(transferBuffer = malloc(bufferSize*sizeof(*transferBuffer)))){
    perror("Failed to allocate transfer buffer");
    exit(EXIT_FAILURE);
}

/* Open Input File Descriptor in Read Only mode */
if((inputFD = open(inputFilename, O_RDONLY | O_SYNC)) < 0){
    perror("Failed to open input file");
    exit(EXIT_FAILURE);
}

/* Open Output File Descriptor in Write Only mode with standard permissions*/
rv = snprintf(outputFilename, MAXFILENAMELENGTH, "%s-%d",
              outputFilenameBase, getpid());
if(rv > MAXFILENAMELENGTH){
    fprintf(stderr, "Output filename length exceeds limit of %d characters.\n",
            MAXFILENAMELENGTH);
    exit(EXIT_FAILURE);
}
else if(rv < 0){
    perror("Failed to generate output filename");
    exit(EXIT_FAILURE);
}
if((outputFD =
    open(outputFilename,
        O_WRONLY | O_CREAT | O_TRUNC | O_SYNC,
        S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP | S_IROTH)) < 0){
    perror("Failed to open output file");
    exit(EXIT_FAILURE);
}

/* Print Status */
fprintf(stdout, "Reading from %s and writing to %s\n",
        inputFilename, outputFilename);

```

```

/* Read from input file and write to output file*/
do{
    /* Read transfersize bytes from input file*/
    bytesRead = read(inputFD, transferBuffer, buffersize);
    if(bytesRead < 0){
        perror("Error reading input file");
        exit(EXIT_FAILURE);
    }
    else{
        totalBytesRead += bytesRead;
        totalReads++;
    }

    /* If all bytes were read, write to output file*/
    if(bytesRead == blocksize){
        bytesWritten = write(outputFD, transferBuffer, bytesRead);
        if(bytesWritten < 0){
            perror("Error writing output file");
            exit(EXIT_FAILURE);
        }
        else{
            totalBytesWritten += bytesWritten;
            totalWrites++;
        }
    }
    /* Otherwise assume we have reached the end of the input file and reset */
    else{
        if(lseek(inputFD, 0, SEEK_SET)){
            perror("Error resetting to beginning of file");
            exit(EXIT_FAILURE);
        }
        inputFileResets++;
    }
}while(totalBytesWritten < transfersize);

/* Output some possibly helpfull info to make it seem like we were doing stuff */
fprintf(stdout, "Read:  %zd bytes in %d reads\n",
        totalBytesRead, totalReads);
fprintf(stdout, "Written: %zd bytes in %d writes\n",
        totalBytesWritten, totalWrites);
fprintf(stdout, "Read input file in %d pass%s\n",
        (inputFileResets + 1), (inputFileResets ? "es" : ""));
fprintf(stdout, "Processed %zd bytes in blocks of %zd bytes\n",
        transfersize, blocksize);

/* Free Buffer */
free(transferBuffer);

/* Close Output File Descriptor */
if(close(outputFD)){
    perror("Failed to close output file");
}

```

```

        exit(EXIT_FAILURE);
    }

    /* Close Input File Descriptor */
    if(close(inputFD)){
        perror("Failed to close input file");
        exit(EXIT_FAILURE);
    }

    /* Calculate pi using statistical methode across all iterations*/
    for(i=0; i<iterations; i++){
        x = (random() % (RADIUS * 2)) - RADIUS;
        y = (random() % (RADIUS * 2)) - RADIUS;
        if(zeroDist(x,y) < RADIUS){
            inCircle++;
        }
        inSquare++;
    }

    /* Finish calculation */
    pCircle = inCircle/inSquare;
    piCalc = pCircle * 4.0;

    /* Print result */
    fprintf(stdout, "pi = %f\n", piCalc);

    return EXIT_SUCCESS;
}

```

```

-pi.c-
/*
 * File: pi.c
 * Author: Andy Sayler
 * Project: CSCI 3753 Programming Assignment 3
 * Create Date: 2012/03/07
 * Modify Date: 2012/03/09
 * Description:
 *     This file contains a simple program for statistically
 *     calculating pi.
 */

```

```

/* Local Includes */
#include <stdlib.h>
#include <stdio.h>
#include <math.h>
#include <errno.h>

```

```

/* Local Defines */
#define DEFAULT_ITERATIONS 1000000
#define RADIUS (RAND_MAX / 2)

```

```

/* Local Functions */

```



```

inline double dist(double x0, double y0, double x1, double y1){
    return sqrt(pow((x1-x0),2) + pow((y1-y0),2));
}

```

```

inline double zeroDist(double x, double y){
    return dist(0, 0, x, y);
}

```

```

int main(int argc, char* argv[]){

    long i;
    long iterations;
    double x, y;
    double inCircle = 0.0;
    double inSquare = 0.0;
    double pCircle = 0.0;
    double piCalc = 0.0;

    /* Process program arguments to select iterations */
    /* Set default iterations if not supplied */
    if(argc < 2){
        iterations = DEFAULT_ITERATIONS;
    }
    /* Set iterations if supplied */
    else{
        iterations = atol(argv[1]);
        if(iterations < 1){
            fprintf(stderr, "Bad iterations value\n");
            exit(EXIT_FAILURE);
        }
    }

    /* Calculate pi using statistical method across all iterations*/
    for(i=0; i<iterations; i++){
        x = (random() % (RADIUS * 2)) - RADIUS;
        y = (random() % (RADIUS * 2)) - RADIUS;
        if(zeroDist(x,y) < RADIUS){
            inCircle++;
        }
        inSquare++;
    }

    /* Finish calculation */
    pCircle = inCircle/inSquare;
    piCalc = pCircle * 4.0;

    /* Print result */
    fprintf(stdout, "pi = %f\n", piCalc);

    return 0;
}

```

```

-rw.c-
/*
 * File: rw.c
 * Author: Andy Saylor
 * Project: CSCI 3753 Programming Assignment 3
 * Create Date: 2012/03/19
 * Modify Date: 2012/03/20
 * Description: A small i/o bound program to copy N bytes from an input
 *              file to an output file. May read the input file multiple
 *              times if N is larger than the size of the input file.
 */

/* Include Flags */
#define _GNU_SOURCE

/* System Includes */
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <errno.h>
#include <sched.h>
#include <fcntl.h>
#include <string.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/wait.h>

/* Local Defines */
#define MAXFILENAMELENGTH 80
#define DEFAULT_INPUTFILENAME "rwinput"
#define DEFAULT_OUTPUTFILENAMEBASE "rwoutput"
#define DEFAULT_BLOCKSIZE 1024
#define DEFAULT_TRANSFERSIZE 1024*100
#define DEFAULT_PROCESSES 1

int main(int argc, char* argv[]){

    int rv;
    int inputFD;
    int outputFD;
    int status;
    int policy;
    long processes;
    struct sched_param param;
    char inputFilename[MAXFILENAMELENGTH];
    char outputFilename[MAXFILENAMELENGTH];
    char outputFilenameBase[MAXFILENAMELENGTH];

    ssize_t transfersize = 0;
    ssize_t blocksize = 0;
    char* transferBuffer = NULL;
    ssize_t buffersize;

```

```

ssize_t bytesRead = 0;
ssize_t totalBytesRead = 0;
int totalReads = 0;
ssize_t bytesWritten = 0;
ssize_t totalBytesWritten = 0;
int totalWrites = 0;
int inputFileResets = 0;

/* Process program arguments to select run-time parameters */
/* Set supplied transfer size or default if not supplied */
if(argc < 2){
    transfersize = DEFAULT_TRANSFERSIZE;
}
else{
    transfersize = atol(argv[1]);
    if(transfersize < 1){
        fprintf(stderr, "Bad transfersize value\n");
        exit(EXIT_FAILURE);
    }
}
/* Set supplied block size or default if not supplied */
if(argc < 3){
    blocksize = DEFAULT_BLOCKSIZE;
}
else{
    blocksize = atol(argv[2]);
    if(blocksize < 1){
        fprintf(stderr, "Bad blocksize value\n");
        exit(EXIT_FAILURE);
    }
}
/* Set supplied input filename or default if not supplied */
if(argc < 4){
    if(strlen(DEFAULT_INPUTFILENAME, MAXFILENAMELENGTH) >=
MAXFILENAMELENGTH){
        fprintf(stderr, "Default input filename too long\n");
        exit(EXIT_FAILURE);
    }
    strncpy(inputFilename, DEFAULT_INPUTFILENAME,
MAXFILENAMELENGTH);
}
else{
    if(strlen(argv[3], MAXFILENAMELENGTH) >= MAXFILENAMELENGTH){
        fprintf(stderr, "Input filename too long\n");
        exit(EXIT_FAILURE);
    }
    strncpy(inputFilename, argv[3], MAXFILENAMELENGTH);
}
/* Set supplied output filename base or default if not supplied */
if(argc < 5){

```

```

        if(strlen(DEFAULT_OUTPUTFILENAMEBASE, MAXFILENAMELENGTH) >=
MAXFILENAMELENGTH){
            fprintf(stderr, "Default output filename base too long\n");
            exit(EXIT_FAILURE);
        }
        strncpy(outputFilenameBase, DEFAULT_OUTPUTFILENAMEBASE,
MAXFILENAMELENGTH);
    }
    else{
        if(strlen(argv[4], MAXFILENAMELENGTH) >= MAXFILENAMELENGTH){
            fprintf(stderr, "Output filename base is too long\n");
            exit(EXIT_FAILURE);
        }
        strncpy(outputFilenameBase, argv[4], MAXFILENAMELENGTH);
    }
    /* Set default policy if not supplied */
    if(argc < 7){
        policy = SCHED_OTHER;
    }
    /* Set policy if supplied */
    if(argc > 6){
        if(!strcmp(argv[6], "SCHED_OTHER")){
            policy = SCHED_OTHER;
        }
        else if(!strcmp(argv[6], "SCHED_FIFO")){
            policy = SCHED_FIFO;
        }
        else if(!strcmp(argv[6], "SCHED_RR")){
            policy = SCHED_RR;
        }
        else{
            fprintf(stderr, "Unhanded scheduling policy\n");
            exit(EXIT_FAILURE);
        }
    }
}
/* Set process to max prioty for given scheduler */
param.sched_priority = sched_get_priority_max(policy);
/* Set new scheduler policy */
fprintf(stdout, "Current Scheduling Policy: %d\n", sched_getscheduler(0));
fprintf(stdout, "Setting Scheduling Policy to: %d\n", policy);
if(sched_setscheduler(0, policy, &param)){
    perror("Error setting scheduler policy");
    exit(EXIT_FAILURE);
}
fprintf(stdout, "New Scheduling Policy: %d\n", sched_getscheduler(0));

/* Modification for forks */
if(argc < 6){
    processes = DEFAULT_PROCESSES;
}

/* Modification for forks */

```

```

if(argc > 5){
    processes = atol(argv[5]);
    if(processes < 1){
        fprintf(stderr, "Bad processes value\n");
        exit(EXIT_FAILURE);
    }
}

/* Modification for forks */
int f;
int g;
for(f=0; f<processes; f++){
    if(fork() == 0){
        execve(argv[1],argv, NULL);
    }
}
for(g=0; g<processes; g++){
    waitpid(0, &status, 0);
}

/* Confirm blocksize is multiple of and less than transfersize*/
if(blocksize > transfersize){
    fprintf(stderr, "blocksize can not exceed transfersize\n");
    exit(EXIT_FAILURE);
}
if(transfersize % blocksize){
    fprintf(stderr, "blocksize must be multiple of transfersize\n");
    exit(EXIT_FAILURE);
}

/* Allocate buffer space */
buffersize = blocksize;
if(!(transferBuffer = malloc(buffersize*sizeof(*transferBuffer)))){
    perror("Failed to allocate transfer buffer");
    exit(EXIT_FAILURE);
}

/* Open Input File Descriptor in Read Only mode */
if((inputFD = open(inputFilename, O_RDONLY | O_SYNC)) < 0){
    perror("Failed to open input file");
    exit(EXIT_FAILURE);
}

/* Open Output File Descriptor in Write Only mode with standard permissions*/
rv = snprintf(outputFilename, MAXFILENAMELENGTH, "%s-%d",
    outputFilenameBase, getpid());
if(rv > MAXFILENAMELENGTH){
    fprintf(stderr, "Output filename length exceeds limit of %d characters.\n",
        MAXFILENAMELENGTH);
    exit(EXIT_FAILURE);
}
else if(rv < 0){

```

```

    perror("Failed to generate output filename");
    exit(EXIT_FAILURE);
}
if((outputFD =
    open(outputFilename,
        O_WRONLY | O_CREAT | O_TRUNC | O_SYNC,
        S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP | S_IROTH)) < 0){
    perror("Failed to open output file");
    exit(EXIT_FAILURE);
}

/* Print Status */
fprintf(stdout, "Reading from %s and writing to %s\n",
    inputFilename, outputFilename);

/* Read from input file and write to output file*/
do{
    /* Read transfersize bytes from input file*/
    bytesRead = read(inputFD, transferBuffer, buffersize);
    if(bytesRead < 0){
        perror("Error reading input file");
        exit(EXIT_FAILURE);
    }
    else{
        totalBytesRead += bytesRead;
        totalReads++;
    }

    /* If all bytes were read, write to output file*/
    if(bytesRead == blocksize){
        bytesWritten = write(outputFD, transferBuffer, bytesRead);
        if(bytesWritten < 0){
            perror("Error writing output file");
            exit(EXIT_FAILURE);
        }
        else{
            totalBytesWritten += bytesWritten;
            totalWrites++;
        }
    }
}
/* Otherwise assume we have reached the end of the input file and reset */
else{
    if(lseek(inputFD, 0, SEEK_SET)){
        perror("Error resetting to beginning of file");
        exit(EXIT_FAILURE);
    }
    inputFileResets++;
}

}while(totalBytesWritten < transfersize);

/* Output some possibly helpfull info to make it seem like we were doing stuff */

```

```

fprintf(stdout, "Read:  %zd bytes in %d reads\n",
          totalBytesRead, totalReads);
fprintf(stdout, "Written: %zd bytes in %d writes\n",
          totalBytesWritten, totalWrites);
fprintf(stdout, "Read input file in %d pass%s\n",
          (inputFileResets + 1), (inputFileResets ? "es" : ""));
fprintf(stdout, "Processed %zd bytes in blocks of %zd bytes\n",
          transfersize, blocksize);

/* Free Buffer */
free(transferBuffer);

/* Close Output File Descriptor */
if(close(outputFD)){
    perror("Failed to close output file");
    exit(EXIT_FAILURE);
}

/* Close Input File Descriptor */
if(close(inputFD)){
    perror("Failed to close input file");
    exit(EXIT_FAILURE);
}

return EXIT_SUCCESS;
}

```

-MakeFile-

```

CC = gcc
CFLAGS = -c -g -Wall -Wextra
LFLAGS = -g -Wall -Wextra

```

```
INPUTFILESIZEMEGABYTES = 1
```

```

KILO = 1024
MEGA = $(shell echo $(KILO)\*$(KILO) | bc)
INPUTFILESIZEBYTES = $(shell echo $(MEGA)\*$(INPUTFILESIZEMEGABYTES) |
bc)
INPUTBLOCKSIZEBYTES = $(KILO)
INPUTBLOCKS = $(shell echo $(INPUTFILESIZEBYTES)/
$(INPUTBLOCKSIZEBYTES) | bc)

```

```
.PHONY: all clean
```

```
all: pi pi-sched rw mixed
```

```

pi: pi.o
    $(CC) $(LFLAGS) $^ -o $@ -lm

```

```
pi-sched: pi-sched.o
```

```

$(CC) $(LFLAGS) $^ -o $@ -lm

rw: rw.o rwinput
    $(CC) $(LFLAGS) rw.o -o $@ -lm

mixed: mixed.o rwinput
    $(CC) $(LFLAGS) mixed.o -o $@ -lm

pi.o: pi.c
    $(CC) $(CFLAGS) $<

pi-sched.o: pi-sched.c
    $(CC) $(CFLAGS) $<

rw.o: rw.c
    $(CC) $(CFLAGS) $<

mixed.o: mixed.c
    $(CC) $(CFLAGS) $<

rwinput: Makefile
    dd if=/dev/urandom of=./rwinput bs=$(INPUTBLOCKSIZEBYTES) count=
$(INPUTBLOCKS)

clean: testclean
    rm -f pi pi-sched rw mixed
    rm -f rwinput
    rm -f *.o
    rm -f *~
    rm -f handout/*~
    rm -f handout/*.log
    rm -f handout/*.aux

testclean:
    rm -f rwoutput*
    rm -f output*

-pi_sched.c-

/*
 * File: pi-sched.c
 * Author: Andy Sayler
 * Project: CSCI 3753 Programming Assignment 3
 * Create Date: 2012/03/07
 * Modify Date: 2012/03/09
 * Description:
 *     This file contains a simple program for statistically
 *     calculating pi using a specific scheduling policy.
 */

/* Local Includes */
#include <stdlib.h>

```



```

#include <stdio.h>
#include <string.h>
#include <math.h>
#include <errno.h>
#include <sched.h>
#include <sys/wait.h>
#include <sys/types.h>
#include <unistd.h>

#define DEFAULT_ITERATIONS 1000000
#define RADIUS (RAND_MAX / 2)
#define DEFAULT_PROCESSES 1

inline double dist(double x0, double y0, double x1, double y1){
    return sqrt(pow((x1-x0),2) + pow((y1-y0),2));
}

inline double zeroDist(double x, double y){
    return dist(0, 0, x, y);
}

int main(int argc, char* argv[]){

    long i;
    long iterations;
    long processes;
    int status;
    struct sched_param param;
    int policy;
    double x, y;
    double inCircle = 0.0;
    double inSquare = 0.0;
    double pCircle = 0.0;
    double piCalc = 0.0;

    /* Process program arguments to select iterations and policy */
    /* Set default iterations if not supplied */
    if(argc < 2){
        iterations = DEFAULT_ITERATIONS;
    }
    /* Set default policy if not supplied */
    if(argc < 3){
        policy = SCHED_OTHER;
    }
    /* Modification for forks */
    if(argc < 4){
        processes = DEFAULT_PROCESSES;
    }
    /* Set iterations if supplied */
    if(argc > 1){
        iterations = atol(argv[1]);
    }

```

```

        if(iterations < 1){
            fprintf(stderr, "Bad iterations value\n");
            exit(EXIT_FAILURE);
        }
    }
    /* Set policy if supplied */
    if(argc > 2){
        if(!strcmp(argv[2], "SCHED_OTHER")){
            policy = SCHED_OTHER;
        }
        else if(!strcmp(argv[2], "SCHED_FIFO")){
            policy = SCHED_FIFO;
        }
        else if(!strcmp(argv[2], "SCHED_RR")){
            policy = SCHED_RR;
        }
        else{
            fprintf(stderr, "Unhanded scheduling policy\n");
            exit(EXIT_FAILURE);
        }
    }
    /* Modification for forks */
    if(argc > 3){
        processes = atol(argv[3]);
        if(processes < 1){
            fprintf(stderr, "Bad processes value\n");
            exit(EXIT_FAILURE);
        }
    }

    /* Set process to max prioty for given scheduler */
    param.sched_priority = sched_get_priority_max(policy);

    /* Set new scheduler policy */
    fprintf(stdout, "Current Scheduling Policy: %d\n", sched_getscheduler(0));
    fprintf(stdout, "Setting Scheduling Policy to: %d\n", policy);
    if(sched_setscheduler(0, policy, &param)){
        perror("Error setting scheduler policy");
        exit(EXIT_FAILURE);
    }
    fprintf(stdout, "New Scheduling Policy: %d\n", sched_getscheduler(0));

    /* Modification for forks */
    int f;
    int g;
    for(f=0; f<processes; f++){
        if(fork() == 0){
            execve(argv[1],argv, NULL);
        }
    }
    for(g=0; g<processes; g++){
        waitpid(0, &status, 0);
    }

```

```

    }

/* Calculate pi using statistical methode across all iterations*/
for(i=0; i<iterations; i++){
    x = (random() % (RADIUS * 2)) - RADIUS;
    y = (random() % (RADIUS * 2)) - RADIUS;
    if(zeroDist(x,y) < RADIUS){
        inCircle++;
    }
    inSquare++;
}

/* Finish calculation */
pCircle = inCircle/inSquare;
piCalc = pCircle * 4.0;

/* Print result */
fprintf(stdout, "pi = %f\n", piCalc);

return 0;
}

```

```

-TestScript-
#!/bin/bash

```

```

#File: testscript
#Author: Andy Saylor
#Project: CSCI 3753 Programming Assignment 3
#Create Date: 2012/03/09
#Modify Date: 2012/03/21
#Description:
#    A simple bash script to run a signle copy of each test case
#    and gather the relevent data.

```

```

ITERATIONS=100000 #1000000000
SML=6
MED=9
LRG=12
BYTESTOCOPY=102400
BLOCKSIZE=1024
INPUTNAME=rwinput
OUTPUTNAME=rwoutput
INPUTFILES=names1.txt names2.txt names3.txt
RESULTS=results.txt
MAKE="make -s"

```

```

echo Building code...
$MAKE clean
$MAKE

```

```

echo Starting test runs...

```

echo Calculating pi over \$ITERATIONS iterations using SCHED\_OTHER with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./pi-sched $ITERATIONS SCHED_OTHER $SML >> output
```

echo Calculating pi over \$ITERATIONS iterations using SCHED\_FIFO with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./pi-sched $ITERATIONS SCHED_FIFO $SML >> output
```

echo Calculating pi over \$ITERATIONS iterations using SCHED\_RR with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./pi-sched $ITERATIONS SCHED_RR $SML >> output
```

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_OTHER with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./rw $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $SML SCHED_OTHER >> output
```

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_FIFO with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./rw $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $SML SCHED_FIFO >> output
```

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_RR with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./rw $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $SML SCHED_RR >> output
```

echo Running mixed CPU/IO bound process

echo using SCHED\_OTHER with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./mixed $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $SML SCHED_OTHER $ITERATIONS >> output
```

echo Running mixed CPU/IO bound process

echo using SCHED\_FIFO with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./mixed $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $SML SCHED_FIFO $ITERATIONS >> output
```

echo Running mixed CPU/IO bound process

echo using SCHED\_RR with 64 simultaneous processes...

```
/usr/bin/time -v sudo ./mixed $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $SML SCHED_RR $ITERATIONS >> output
```

echo Calculating pi over \$ITERATIONS iterations using SCHED\_OTHER with 512 simultaneous processes...

```
/usr/bin/time -v sudo ./pi-sched $ITERATIONS SCHED_OTHER $MED >> output
```

echo Calculating pi over \$ITERATIONS iterations using SCHED\_FIFO with 512 simultaneous processes...

```
/usr/bin/time -v sudo ./pi-sched $ITERATIONS SCHED_FIFO $MED >> output
```

echo Calculating pi over \$ITERATIONS iterations using SCHED\_RR with 512 simultaneous processes...

/usr/bin/time -v sudo ./pi-sched \$ITERATIONS SCHED\_RR \$MED >> output

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_OTHER with 512 simultaneous processes...

/usr/bin/time -v sudo ./rw \$BYTESTOCOPY \$BLOCKSIZE \$INPUTNAME  
\$OUTPUTNAME \$MED SCHED\_OTHER >> output

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_FIFO with 512 simultaneous processes...

/usr/bin/time -v sudo ./rw \$BYTESTOCOPY \$BLOCKSIZE \$INPUTNAME  
\$OUTPUTNAME \$MED SCHED\_FIFO >> output

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_RR with 512 simultaneous processes...

/usr/bin/time -v sudo ./rw \$BYTESTOCOPY \$BLOCKSIZE \$INPUTNAME  
\$OUTPUTNAME \$MED SCHED\_RR >> output

echo Running mixed CPU/IO bound process

echo using SCHED\_OTHER with 512 simultaneous processes...

/usr/bin/time -v sudo ./mixed \$BYTESTOCOPY \$BLOCKSIZE \$INPUTNAME  
\$OUTPUTNAME \$MED SCHED\_OTHER \$ITERATIONS >> output

echo Running mixed CPU/IO bound process

echo using SCHED\_FIFO with 512 simultaneous processes...

/usr/bin/time -v sudo ./mixed \$BYTESTOCOPY \$BLOCKSIZE \$INPUTNAME  
\$OUTPUTNAME \$MED SCHED\_FIFO \$ITERATIONS >> output

echo Running mixed CPU/IO bound process

echo using SCHED\_RR with 512 simultaneous processes...

/usr/bin/time -v sudo ./mixed \$BYTESTOCOPY \$BLOCKSIZE \$INPUTNAME  
\$OUTPUTNAME \$MED SCHED\_RR \$ITERATIONS >> output

echo Calculating pi over \$ITERATIONS iterations using SCHED\_OTHER with 4096  
simultaneous processes...

/usr/bin/time -v sudo ./pi-sched \$ITERATIONS SCHED\_OTHER \$LRG >> output

echo Calculating pi over \$ITERATIONS iterations using SCHED\_FIFO with 4096  
simultaneous processes...

/usr/bin/time -v sudo ./pi-sched \$ITERATIONS SCHED\_FIFO \$LRG >> output

echo Calculating pi over \$ITERATIONS iterations using SCHED\_RR with 4096  
simultaneous processes...

/usr/bin/time -v sudo ./pi-sched \$ITERATIONS SCHED\_RR \$LRG >> output

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_OTHER with 4096 simultaneous processes...

/usr/bin/time -v sudo ./rw \$BYTESTOCOPY \$BLOCKSIZE \$INPUTNAME  
\$OUTPUTNAME \$LRG SCHED\_OTHER >> output

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_FIFO with 4096 simultaneous processes...

```
/usr/bin/time -v sudo ./rw $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $LRG SCHED_FIFO >> output
```

echo Copying \$BYTESTOCOPY bytes in blocks of \$BLOCKSIZE from rwinput to rwoutput  
echo using SCHED\_RR with 4096 simultaneous processes...

```
/usr/bin/time -v sudo ./rw $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $LRG SCHED_RR >> output
```

echo Running mixed CPU/IO bound process

echo using SCHED\_OTHER with 4096 simultaneous processes...

```
/usr/bin/time -v sudo ./mixed $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $LRG SCHED_OTHER $ITERATIONS >> output
```

echo Running mixed CPU/IO bound process

echo using SCHED\_FIFO with 4096 simultaneous processes...

```
/usr/bin/time -v sudo ./mixed $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $LRG SCHED_FIFO $ITERATIONS >> output
```

echo Running mixed CPU/IO bound process

echo using SCHED\_RR with 4096 simultaneous processes...

```
/usr/bin/time -v sudo ./mixed $BYTESTOCOPY $BLOCKSIZE $INPUTNAME  
$OUTPUTNAME $LRG SCHED_RR $ITERATIONS >> output
```

-util.c-

/\*

\* File: util.c

\* Author: Andy Saylor

\* Project: CSCI 3753 Programming Assignment 2

\* Create Date: 2012/02/01

\* Modify Date: 2012/02/01

\* Description:

\* This file contains declarations of utility functions for  
\* Programming Assignment 2.

\*

\*/

#include "util.h"

```
int dnslookup(const char* hostname, char* firstIPstr, int maxSize){
```

```
/* Local vars */
```

```
struct addrinfo* headresult = NULL;
```

```
struct addrinfo* result = NULL;
```

```
struct sockaddr_in* ipv4sock = NULL;
```

```
struct sockaddr_in6* ipv6sock = NULL;
```

```
struct in_addr* ipv4addr = NULL;
```

```
struct in6_addr* ipv6addr = NULL;
```

```
char ipv4str[INET_ADDRSTRLEN];
```

```
char ipv6str[INET6_ADDRSTRLEN];
```

```
char ipstr[INET6_ADDRSTRLEN];
```

```
int addrError = 0;
```

```

    /* DEBUG: Print Hostname */
#ifdef UTIL_DEBUG
    fprintf(stderr, "%s\n", hostname);
#endif

    /* Lookup Hostname */
    addrError = getaddrinfo(hostname, NULL, NULL, &headresult);
    if(addrError){
        fprintf(stderr, "Error looking up Address: %s\n",
            gai_strerror(addrError));
        return UTIL_FAILURE;
    }
    /* Loop Through result Linked List */
    for(result=headresult; result != NULL; result = result->ai_next){
        /* Extract IP Address and Convert to String */
        if(result->ai_addr->sa_family == AF_INET){
            /* IPv4 Address Handling */
            ipv4sock = (struct sockaddr_in*)(result->ai_addr);
            ipv4addr = &(ipv4sock->sin_addr);
            if(!inet_ntop(result->ai_family, ipv4addr,
                ipv4str, sizeof(ipv4str))){
                perror("Error Converting IP to String");
                return UTIL_FAILURE;
            }
#ifdef UTIL_DEBUG
            fprintf(stdout, "%s\n", ipv4str);
#endif
            strncpy(ipstr, ipv4str, sizeof(ipstr));
            ipstr[sizeof(ipstr)-1] = '\0';
        }
        else if(result->ai_addr->sa_family == AF_INET6){
            /* IPv6 Handling */
            ipv6sock = (struct sockaddr_in6*)(result->ai_addr);
            ipv6addr = &(ipv6sock->sin6_addr);
            if(!inet_ntop(result->ai_family, ipv6addr,
                ipv6str, sizeof(ipv6str))){
                perror("Error Converting IP to String");
                return UTIL_FAILURE;
            }
#ifdef UTIL_DEBUG
            fprintf(stdout, "IPv6 Address: Not Handled\n");
#endif
            strncpy(ipstr, ipv6str, sizeof(ipstr));
            ipstr[sizeof(ipstr)-1] = '\0';
        }
        else{
            /* Unhandlded Protocol Handling */
#ifdef UTIL_DEBUG
            fprintf(stdout, "Unknown Protocol: Not Handled\n");
#endif
            strncpy(ipstr, "UNHANDELED", sizeof(ipstr));
            ipstr[sizeof(ipstr)-1] = '\0';
        }
    }

```

```

        }
        /* Save First IP Address */
        if(result==headresult){
            strncpy(firstIPstr, ipstr, maxSize);
            firstIPstr[maxSize-1] = '\0';
        }
    }

    /* Cleanup */
    freeaddrinfo(headresult);

    return UTIL_SUCCESS;
}

-util.h-
/*
 * File: util.h
 * Author: Andy Saylor
 * Project: CSCI 3753 Programming Assignment 2
 * Create Date: 2012/02/01
 * Modify Date: 2012/02/01
 * Description:
 *     This file contains declarations of utility functions for
 *     Programming Assignment 2.
 */

#ifndef UTIL_H
#define UTIL_H

/* Define the following to enable debug statments */
// #define UTIL_DEBUG

#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <errno.h>

#include <arpa/inet.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

#define UTIL_FAILURE -1
#define UTIL_SUCCESS 0
#define NUM_ADDR 50

/* Fuction to return the first IP address found
 * for hostname. IP address returned as string
 * firstIPstr of size maxsize
 */
int dnslookup(const char* hostname,

```



```
char* firstIPstr,  
int maxSize);
```